

# PyTorch for Beginners: Learn the Basics of Deep Learning with Python

---

## Course Content

1. Introduction to Classes in python
2. introduction to PyTorch
  - a. What is PyTorch?
  - b. Why choose PyTorch for deep learning?
  - c. Installation and setup
3. Tensors and Operations
  - a. Tensors: The fundamental data structure in PyTorch
  - b. Creating tensors: **torch.Tensor**, **torch.zeros**, **torch.ones**, **torch.rand**
  - c. Basic operations: elementwise addition, multiplication, etc.
  - d. Indexing and slicing tensors
  - e. Broadcasting
4. Autograd: Automatic Differentiation
  - a. Automatic differentiation in PyTorch
  - b. Creating tensors with gradients: **torch.tensor** with **requires\_grad=True**
  - c. Calculating gradients with **.backward()**
  - d. Gradient accumulation and clearing gradients
  - e. Using **.grad** and **.detach()**
5. Building and Training a Neural Network
  - a. Defining a simple neural network using **nn.Module**
  - b. Specifying layers and activation functions
  - c. Forward and backward passes
  - d. Defining loss functions and optimizers
  - e. Training loops
  - f. Using the **nn.Sequential** container
6. Loading and Preprocessing Data
  - a. Loading custom datasets with **torch.utils.data.Dataset**
  - b. Data transformations using **torchvision.transforms**
  - c. Creating data loaders with **torch.utils.data.DataLoader**
  - d. Using the **torchvision** library for common datasets
7. Model Evaluation and Validation
  - a. Making predictions with trained models
  - b. Computing accuracy and other metrics
  - c. Validation and overfitting
8. Saving and Loading Models

- a. Saving and loading model weights
  - b. Serialisation with **torch.save** and **torch.load**
  - c. Saving and loading entire models
- 9. Transfer Learning and FineTuning
  - a. Using pretrained models from **torchvision.models**
  - b. Feature extraction and finetuning
  - c. Modifying the architecture for transfer learning
- 10. GPU Acceleration
  - a. Utilising GPUs for faster training
  - b. Moving tensors and models to the GPU
  - c. GPUrelated functions: **cuda()**, **cpu()**, **to()**
- 11. Customizing Training and Visualizing Results
  - a. Adding custom loss functions and metrics
  - b. TensorBoard integration for visualisation
  - c. Plotting training curves using matplotlib

# Introduction to PyTorch and its Significance in Deep Learning

PyTorch is an open source machine learning framework that provides a flexible and dynamic environment for building and training neural networks. It has gained immense popularity within the deep learning community due to its intuitive design, dynamic computation graph, and seamless GPU acceleration. Developed by Facebooks AI Research lab (FAIR), PyTorch offers several key features that make it stand out:

## 1. Dynamic Computation Graph:

PyTorchs dynamic computation graph allows you to create complex neural network architectures on the fly. Unlike static computation graphs in some other frameworks, PyTorchs dynamic graph enables more flexible and intuitive model building. This is particularly beneficial when dealing with variable length inputs or dynamic architectures.

## 2. Eager Execution:

With PyTorch, you can execute operations immediately as they are defined. This enables real time debugging, experimentation, and iterative model development. It simplifies the process of understanding and diagnosing issues in your code, making the learning curve smoother for newcomers.

### 3. GPU Acceleration:

PyTorch seamlessly supports GPU acceleration through CUDA, which significantly speeds up the training and evaluation of deep learning models. This feature is crucial for handling large datasets and complex neural network architectures efficiently.

### 4. Strong Community and ResearchFriendly:

PyTorch is widely used in academia and research due to its ease of use and a rich ecosystem of libraries. Its flexible nature makes it a natural choice for experimenting with novel ideas and quickly implementing cutting edge research papers.

### 5. Neural Network Debugging:

PyTorch's dynamic nature simplifies debugging by allowing you to examine and modify intermediate variables during computation. This feature assists in identifying errors and understanding how different components of your model interact.

### 6. Extensive Libraries:

PyTorch provides a rich collection of libraries for tasks such as computer vision (TorchVision), natural language processing (TorchText), and more. These libraries facilitate various aspects of deep learning pipeline development.

### 7. Growing Adoption:

Due to its simplicity and power, PyTorch has experienced rapid adoption by both researchers and practitioners. It has become a go-to framework for various machine learning competitions and projects.

In summary, PyTorch's dynamic and intuitive design, strong community support, and compatibility with GPU acceleration make it a highly effective tool for deep learning research, experimentation, and development. Its userfriendly interface and seamless integration with Python have contributed to its growing significance in the ever evolving field of deep learning.

## Tensors in PyTorch: Types and Operations

Tensors are the fundamental data structures in PyTorch that represent multidimensional arrays. They are similar to NumPy arrays but with additional capabilities for GPU

acceleration and automatic differentiation. Tensors come in various types and support a wide range of operations for mathematical and deep learning computations.

## Tensor Types

PyTorch supports several tensor types, each suited for different use cases:

1. `torch.FloatTensor`: The default tensor type, representing a multidimensional array of singleprecision floatingpoint numbers.
2. `torch.DoubleTensor`: A tensor type with doubleprecision floatingpoint numbers.
3. `torch.IntTensor`, `torch.LongTensor`: Tensors with integer data types.
4. `torch.ByteTensor`: A binary (0/1) tensor used for logical operations.
5. `torch.CharTensor`, `torch.ShortTensor`: Tensors with smaller integer data types.

Further Reading:

[https://pytorch.org/tutorials/beginner/basics/tensorqs\\_tutorial.html](https://pytorch.org/tutorials/beginner/basics/tensorqs_tutorial.html)

## Automatic Differentiation in PyTorch

Automatic differentiation is a crucial concept in machine learning and deep learning. It refers to the process of automatically computing the derivatives of functions with respect to their input variables. This is particularly valuable when training neural networks, as gradient information is essential for optimization algorithms like gradient descent.

In PyTorch, the `autograd` package provides automatic differentiation functionality. It keeps track of operations performed on tensors and their gradients, enabling you to compute gradients of arbitrary functions without manually deriving them. This greatly simplifies the process of calculating gradients, making it easier to train complex neural networks.

## Calculating Gradients using Autograd

Lets illustrate how to calculate gradients using PyTorchs autograd package:

```
import torch

# Create a tensor with requires_grad=True
x = torch.tensor(2.0, requires_grad=True)

# Define a function
y = x2 + 3*x + 1

# Calculate gradients
y.backward()

# Print the gradient
print("Gradient of y with respect to x:", x.grad)
```

In this example, we create a tensor `x` with `requires_grad=True`, indicating that we want to track gradients with respect to `x`. We then define a simple function `y` that depends on `x`. By calling `y.backward()`, PyTorch calculates the gradient of `y` with respect to `x`. The result is accessed using the `.grad` attribute of the input tensor.

Output

**Gradient of y with respect to x: 7.0**

Here, the calculated gradient of the function  $y = x^2 + 3x + 1$  with respect to `x` is 7.0, which is consistent with manual differentiation.

Automatic differentiation in PyTorch allows you to focus on designing and implementing your neural networks architecture and loss functions without having to manually compute gradients. This greatly streamlines the training process and encourages experimentation with different model architectures.

Further Reading:

[https://pytorch.org/tutorials/beginner/basics/autogradqs\\_tutorial.html](https://pytorch.org/tutorials/beginner/basics/autogradqs_tutorial.html)

## Building and Training a Convolutional Neural Network (CNN)

Convolutional Neural Networks (CNNs) are widely used for image classification, object detection, and other computer vision tasks. Heres a stepbystep guide on how to build and train a CNN using PyTorch:

## 1. Import Necessary Libraries

```
import torch
import torch.nn as nn
import torch.optim as optim
```

## 2. Define the CNN Architecture

```
class CNNModel(nn.Module):
    def __init__(self, num_classes):
        super(CNNModel, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3,
padding=1)
        self.relu = nn.ReLU()
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(16 * 112 * 112, 128)
        self.fc2 = nn.Linear(128, num_classes)

    def forward(self, x):
        # print('input', x.shape)
        x = self.pool(self.relu(self.conv1(x)))
        # print('pool', x.shape)
        x = x.view(-1, 16 * 112 * 112)
        # print('view', x.shape)
        x = self.relu(self.fc1(x))
        # print('relu', x.shape)
        x = self.fc2(x)
        # print('output', x.shape)
        return x
```

Define your CNN architecture by creating a class that inherits from nn.Module. This class will contain the layers of your network.

```
# Create an instance of the model
model = CNNModel(num_classes)
```

### 3. Define Loss Function and Optimizer

```
# Initialize loss and optimizer  
optimizer = optim.Adam(model.parameters(), lr=learning_rate)  
criterion = nn.CrossEntropyLoss()
```

### 4. Load and Preprocess Data

Load your dataset using torchvision and apply necessary transformations.

```
import torchvision  
import torchvision.transforms as transforms  
  
transform = transforms.Compose(  
    [transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])  
  
trainset = torchvision.datasets.CIFAR10(root=../data, train=True, download=True,  
    transform=transform)  
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True)
```

### 5. Training Loop

```
for epoch in range(epochs): # Number of training epochs  
    running_loss = 0.0  
    for i, data in enumerate(trainloader, 0):  
        inputs, labels = data  
  
        optimizer.zero_grad()  
  
        outputs = cnn(inputs)  
        loss = criterion(outputs, labels)  
        loss.backward()  
        optimizer.step()  
  
        running_loss += loss.item()  
    print(f"Epoch {epoch + 1}, Loss: {running_loss / len(trainloader)}")
```

### 6. Evaluating the Model

You can use a similar loop to evaluate the models performance on a validation or test dataset.

```
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        inputs, labels = data
        outputs = cnn(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f"Accuracy of the network on the test images: {100 * correct / total}%")
```

This is a basic outline of building and training a CNN using PyTorch. You can further enhance the model by adding more layers, applying regularization techniques, and experimenting with different hyperparameters.

## Loss Functions and Optimizers

### Loss Functions

A loss function, also known as a cost function or objective function, quantifies how well a machine learning models predictions match the actual ground truth labels in the training data. The goal is to minimize this loss during training. Different tasks, such as classification, regression, and more complex problems, require different types of loss functions. For example, the mean squared error (MSE) loss is often used for regression tasks, while the crossentropy loss is commonly used for classification tasks. The choice of the appropriate loss function is crucial as it guides the learning process of the model.

#### 1. Mean Squared Error (MSE) Loss:

`nn.MSELoss()`

Used for regression tasks.

Measures the average squared difference between predicted and target values.

#### 2. CrossEntropy Loss:

`nn.CrossEntropyLoss()`

Commonly used for classification tasks.

Combines softmax activation and negative loglikelihood loss.

Suitable for multiclass classification problems.



### 3. Binary CrossEntropy Loss (BCELoss):

`nn.BCELoss()`

Used for binary classification tasks.

Measures the difference between predicted and target values for each class.

### 4. Smooth L1 Loss:

`nn.SmoothL1Loss()`

A variant of the L1 loss that is less sensitive to outliers.

Commonly used in object detection.

### 5. Huber Loss:

`nn.SmoothL1Loss(delta=1.0)`

Similar to Smooth L1 Loss, with the ability to transition between L1 and L2 loss.

### 6. KullbackLeibler Divergence (KLDivLoss):

`nn.KLDivLoss()`

Measures the difference between two probability distributions.

### 7. Hinge Embedding Loss:

`nn.HingeEmbeddingLoss()`

Used for training classifiers in marginbased ranking problems.

### Further Reading:

<https://pytorch.org/docs/stable/nn.html#loss-functions>

## Optimizers

Optimizers are algorithms that adjust the parameters of a model based on the gradients of the loss function with respect to those parameters. The optimization process aims to find the parameter values that minimize the loss function and, consequently, improve the models performance. Common optimization algorithms include stochastic gradient descent (SGD), Adam, RMSprop, and more. These algorithms use techniques like gradient descent and momentum to iteratively update the models parameters. The choice of optimizer and its hyperparameters can impact the convergence speed and final performance of the model during training.

### 1. Stochastic Gradient Descent (SGD):

`optim.SGD(params, lr=0.01, momentum=0, dampening=0, weight_decay=0, nesterov=False)`

## 2. Adam Optimizer:

```
optim.Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e08, weight_decay=0, amsgrad=False)
```

## 3. RMSprop Optimizer:

```
optim.RMSprop(params, lr=0.01, alpha=0.99, eps=1e08, weight_decay=0, momentum=0, centered=False)
```

## 4. Adagrad Optimizer:

```
optim.Adagrad(params, lr=0.01, lr_decay=0, weight_decay=0, initial_accumulator_value=0)
```

## 5. Adadelta Optimizer:

```
optim.Adadelta(params, lr=1.0, rho=0.9, eps=1e06, weight_decay=0)
```

## 6. Sparse Adam Optimizer:

```
optim.SparseAdam(params, lr=0.001, betas=(0.9, 0.999), eps=1e08)
```

## 7. LBFGS Optimizer:

```
optim.LBFGS(params, lr=1, max_iter=20, max_eval=None, tolerance_grad=1e07, tolerance_change=1e09, history_size=100, line_search_fn=None)
```

## Further Reading:

<https://pytorch.org/docs/stable/optim.html>

## Training Loop

The training loop is the core process of training a machine learning model. It iterates through the dataset multiple times, updating the models parameters to minimize the loss function. The training loop typically involves the following steps:

1. **Forward Pass:** Pass a batch of input data through the model to make predictions.
2. **Loss Computation:** Calculate the loss between the models predictions and the actual labels using the chosen loss function.
3. **Backward Pass (Backpropagation):** Compute the gradients of the loss with respect to the models parameters using automatic differentiation (autograd in PyTorch).
4. **Optimizer Step:** Update the models parameters using the gradients computed in the backward pass and the chosen optimizer algorithm.

5. **Iteration** : Repeat the process for a predefined number of epochs (complete passes through the dataset).

During the training loop, the model's performance gradually improves as the loss decreases. The loop helps the model learn the patterns in the data and adjust its parameters to make better predictions. Regularization techniques, such as weight decay and dropout, can also be applied within the training loop to prevent overfitting.

In summary, loss functions quantify the model's performance, optimizers adjust the model's parameters, and the training loop iteratively improves the model's predictions based on the data and optimization techniques. These components together drive the learning process of a machine learning model.

In PyTorch, there are several built-in loss functions and optimizers that you can use for training your machine learning models. Here's an overview of some commonly used loss functions and optimizers:

## Loading and Preprocessing Data using torchvision

torchvision is a PyTorch library that provides a wide range of tools for working with image data, including loading datasets, applying transformations, and creating data loaders. It simplifies the process of preparing image data for training deep learning models. Let's go through the steps of loading and preprocessing data using torchvision.

### 1. Import Necessary Libraries:

```
import torchvision
import torchvision.transforms as transforms
```

### 2. Define Data Transformations:

Transforms are used to preprocess and augment your data. You can define a series of transformations to be applied to each image in the dataset. Common transformations include resizing, normalizing, and data augmentation (e.g., random flips, rotations).

```
transform = transforms.Compose([
    transforms.Resize((128, 128)), # Resize images to 128x128 pixels
    transforms.RandomHorizontalFlip(), # Randomly flip images horizontally
    transforms.ToTensor(), # Convert images to PyTorch tensors
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]) # Normalize pixel values
])
```

### 3. Load Dataset and Apply Transformations:

Use `torchvision.datasets` to load datasets like CIFAR-10, MNIST, ImageNet, etc. Apply the defined transformations to the dataset.

```
trainset = torchvision.datasets.CIFAR10(root=../data, train=True, download=True,  
transform=transform)
```

### 4. Create Data Loaders:

Data loaders are responsible for creating batches of data during training. They shuffle and batch the data, making it ready for model training.

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size=32, shuffle=True)
```

### 5. Iterating Through Data Batches:

You can now iterate through the data loader to obtain batches of preprocessed images and labels during training.

```
for images, labels in trainloader:  
    # Your training code here
```

In this example, we've loaded the CIFAR-10 dataset, applied a series of transformations, and created a data loader to iterate through batches of data. The `transforms.Compose` function allows you to chain multiple transformations together, ensuring consistent preprocessing across all images.

By using `torchvision`, you can significantly simplify the process of loading, preprocessing, and augmenting image data for training your deep learning models. This helps ensure that your data is properly prepared and standardized before being fed into the model.

Further Reading:

[https://pytorch.org/tutorials/beginner/basics/data\\_tutorial.html](https://pytorch.org/tutorials/beginner/basics/data_tutorial.html)

## Role of Transforms in PyTorch

In the context of PyTorch's `torchvision` library, a transform is a set of data preprocessing or augmentation operations applied to input data before it is used for training, validation, or testing. Transforms are an essential component of the data pipeline in machine learning and computer vision tasks. Their primary role is to prepare and preprocess data in a consistent

and standardized manner, which can lead to more effective model training and improved results.

Here's why transforms are necessary and what roles they play:

## 1. Data Preprocessing:

Transforms help preprocess raw data into a format that is suitable for training deep learning models. This includes tasks like resizing images to a uniform size, normalizing pixel values to a specific range, and converting data into the appropriate data types.

## 2. Standardization:

By applying the same set of transformations to all data samples, you ensure that the input data is standardized. This helps in removing biases or variations that might be present in the raw data, leading to more consistent and stable model performance.

## 3. Data Augmentation:

Data augmentation involves applying a variety of image transformations, such as rotations, flips, and cropping, to increase the diversity of the training dataset. Augmentation helps the model become more robust by exposing it to different variations of the same image, reducing overfitting and improving generalization.

## 4. Handling Image Channels:

For RGB images, each pixel has three color channels (red, green, blue). Transforms can rearrange the channel order or apply color space adjustments to ensure that the input data is correctly formatted for the models input requirements.

## 5. Normalization:

Transforms often include normalization steps to bring pixel values within a certain range. This helps stabilize the learning process, as neural networks tend to work better with inputs that have similar scales. Commonly, pixel values are normalized to the range  $[0, 1]$  or standardized with a mean of 0 and a standard deviation of 1.

## 6. Tensor Conversion:

Transforms also convert data into PyTorch tensors, which are the primary data type used by neural networks. This conversion enables seamless integration of the data with PyTorchs computation graph and GPU acceleration.

In summary, transforms play a critical role in preparing, standardizing, and augmenting data for training deep learning models. They ensure that the input data is consistent, properly formatted, and suitable for the specific requirements of the neural network architecture. By applying transformations to your data, you can improve the models ability to learn and generalize patterns from the training data to unseen examples.

## Preloaded Dataset in Torchvision library

The torchvision library provides preloaded datasets commonly used for training and evaluating machine learning models, particularly in the context of computer vision. These datasets are available for quick experimentation and testing. Here are some of the most commonly used preloaded datasets in torchvision, along with explanations of how to use them:

### 1. CIFAR-10 and CIFAR-100:

The CIFAR-10 dataset contains 60,000 32x32 color images in 10 different classes, with 6,000 images per class. The CIFAR-100 dataset is similar but has 100 classes with 600 images each.

Usage:

```
import torchvision
import torchvision.transforms as transforms

# Define transformations
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, 0.5,
0.5), (0.5, 0.5, 0.5))])

# Load CIFAR-10 dataset
trainset = torchvision.datasets.CIFAR10(root= './data', train=True, download=True,
transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=32, shuffle=True)
```

### 2. MNIST:

The MNIST dataset contains 28x28 grayscale images of handwritten digits from 0 to 9.

Usage:

```
python
import torchvision
```

```
import torchvision.transforms as transforms

# Define transformations
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),
(0.5,))])

# Load MNIST dataset
trainset = torchvision.datasets.MNIST(root=../data, train=True, download=True,
transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=32, shuffle=True)
```

### 3. ImageNet:

The ImageNet dataset is a large-scale dataset with millions of labeled images across thousands of classes. torchvision provides a way to download and work with subsets of the ImageNet dataset.

Usage:

```
python
import torchvision
import torchvision.transforms as transforms

# Define transformations
transform = transforms.Compose([transforms.Resize((224, 224)), transforms.ToTensor(), ...])

# Load ImageNet dataset (subset)
trainset = torchvision.datasets.ImageNet(root=../data, split=train, download=True,
transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=32, shuffle=True)
```

These are just a few examples of preloaded datasets provided by torchvision. Each dataset has specific attributes and classes for easy integration with your machine learning models. By using these datasets, you can quickly experiment with various model architectures, loss functions, and optimizers without the need to manually prepare and preprocess the data.

## Data augmentation using “torchvision.transforms”

transforms in PyTorchs torchvision library provide a variety of parameters that you can use to augment and preprocess your data. Data augmentation helps in increasing the diversity of your training dataset, leading to better generalization of your model. Here are some common parameters that can be used in transforms for data augmentation:

1. RandomHorizontalFlip:
  - transforms.RandomHorizontalFlip(p=0.5): Flips the image horizontally with a probability p.
2. RandomVerticalFlip:
  - transforms.RandomVerticalFlip(p=0.5): Flips the image vertically with a probability p.
3. RandomRotation:
  - transforms.RandomRotation(degrees, resample=False, expand=False): Randomly rotates the image by a specified number of degrees.
4. ColorJitter:
  - transforms.ColorJitter(brightness=0, contrast=0, saturation=0, hue=0): Adjusts brightness, contrast, saturation, and hue of the image.
5. RandomAffine:
  - transforms.RandomAffine(degrees, translate=None, scale=None, shear=None): Applies random affine transformations to the image.
6. RandomCrop:
  - transforms.RandomCrop(size, padding=None, pad\_if\_needed=False, fill=0, padding\_mode=constant): Randomly crops the image to a specified size.
7. RandomResizedCrop:
  - transforms.RandomResizedCrop(size, scale=(0.08, 1.0), ratio=(0.75, 1.3333)): Randomly crops and resizes the image while maintaining aspect ratio.
8. GaussianBlur:
  - transforms.GaussianBlur(kernel\_size, sigma=(0.1, 2.0)): Applies Gaussian blur to the image.
9. Lambda:
  - transforms.Lambda(lambda\_func): Applies a custom lambda function to the image.
10. Normalize:
  - transforms.Normalize(mean, std): Normalizes pixel values of the image with specified mean and standard deviation.
11. ToTensor:
  - transforms.ToTensor(): Converts the image to a PyTorch tensor.

These are just a few examples of the available transformations and their parameters in torchvision. You can combine multiple transformations using transforms.Compose to create a pipeline for augmenting and preprocessing your data before its fed into your model for training.



# Evaluating and Validating the Trained Model

Evaluating and validating a trained model is a crucial step to assess its performance on unseen data and ensure its generalization capability. This is typically done using a validation dataset that the model hasn't seen during training. Metrics such as accuracy, precision, recall, F1-score, and confusion matrix are commonly used to evaluate the model's performance. Here's how to evaluate a trained model and calculate accuracy using PyTorch:

## 1. Load and Prepare Validation Data:

Load the validation dataset using the same preprocessing steps used for the training data.

## 2. Set the Model to Evaluation Mode:

Before evaluating the model, set it to evaluation mode using `model.eval()`. This turns off certain layers like dropout and batch normalization, ensuring consistent behavior during inference.

```
python
model.eval()
```

## 3. Evaluate the Model:

Iterate through the validation dataset, forward-pass the data through the model, and compare the predicted labels with the ground truth labels to calculate metrics.

```
python
correct = 0
total = 0
with torch.no_grad(): # Disable gradient calculation
    for images, labels in val_loader:
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```

```
accuracy = 100 * correct / total
print(f"Accuracy: {accuracy:.2f}%")
```

## 4. Calculate Other Metrics:

You can calculate other metrics like precision, recall, F1-score, and the confusion matrix using tools like scikit-learn or directly using PyTorch.

```
python
from sklearn.metrics import classification_report, confusion_matrix
```

```
# Get predictions and true labels
```

```

all_predictions = []
all_labels = []
with torch.no_grad():
    for images, labels in val_loader:
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        all_predictions.extend(predicted.tolist())
        all_labels.extend(labels.tolist())

# Calculate metrics
report = classification_report(all_labels, all_predictions, target_names=class_names)
conf_matrix = confusion_matrix(all_labels, all_predictions)
print("Classification Report:\n", report)
print("Confusion Matrix:\n", conf_matrix)

```

Remember to replace `val\_loader` with the appropriate validation data loader and `class\_names` with your datasets class names.

Evaluating and validating the model ensures that it performs well on unseen data and helps you identify any issues with overfitting or underfitting. A combination of accuracy and other relevant metrics provides a comprehensive understanding of your models performance.

## Saving and Loading Model Weights

After training a machine learning model, its important to save its learned parameters (weights) so that you can use the model for inference or resume training later without starting from scratch. Saving and loading model weights is a common practice in deep learning. PyTorch provides a simple way to save and load model weights using the `torch.save()` and `torch.load()` functions.

### 1. Saving Model Weights:

You can save the models state dictionary, which includes all the learned parameters and information about the models architecture.

```

python
# Save the models state dictionary
torch.save(model.state_dict(), model_weights.pth)

```

### 2. Loading Model Weights:

To load the saved model weights into a model, you need to instantiate the model and load the saved state dictionary.

```
python
# Create an instance of the model
model = YourModelClass()

# Load the saved state dictionary
model.load_state_dict(torch.load(model_weights.pth))
```

Heres a step-by-step demonstration of how to save and load model weights:

### 1. Saving Model Weights:

```
python
import torch
import torch.nn as nn

# Define a simple model
class SimpleModel(nn.Module):
    def __init__(self):
        super(SimpleModel, self).__init__()
        self.fc = nn.Linear(10, 5)

model = SimpleModel()

# Save model weights
torch.save(model.state_dict(), model_weights.pth)
```

### 2. Loading Model Weights:

```
python
import torch
import torch.nn as nn

# Define the same simple model
class SimpleModel(nn.Module):
    def __init__(self):
        super(SimpleModel, self).__init__()
        self.fc = nn.Linear(10, 5)

# Instantiate the model
model = SimpleModel()

# Load saved weights
model.load_state_dict(torch.load(model_weights.pth))
model.eval() # Set to evaluation mode
```

In this example, we first define a simple model, save its weights, and then load the saved weights into another instance of the same model. After loading the weights, remember to set the model to evaluation mode if you're using it for inference.

Saving and loading model weights is essential for sharing, deploying, and reusing trained models without the need to retrain them from scratch. It's especially important when working with large and complex architectures that require significant training time.

## Saving Model Checkpoints During Training

Saving model checkpoints during training is a common practice to keep track of the model's progress and allow resumption from specific points in training. It involves saving not only the model's weights but also other important training parameters such as optimizer state, current epoch, and more. This allows you to resume training from where you left off in case of interruptions or when you want to fine-tune a model. Here's how to save and use model checkpoints:

### 1. Save Model Checkpoint:

During training, save the model's state dictionary, optimizer state, and other relevant training parameters.

```
python
import torch
import torch.nn as nn
import torch.optim as optim

# Define your model and optimizer
model = ...
optimizer = optim.SGD(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

# Training loop
num_epochs = 10
for epoch in range(num_epochs):
    # Training code here
    # ...

    # Save model checkpoint
    checkpoint = {
        epoch: epoch + 1,
        state_dict: model.state_dict(),
        optimizer: optimizer.state_dict(),
        loss: loss,
        # Other training parameters
    }
    torch.save(checkpoint, f'checkpoint_epoch_{epoch + 1}.pth')
```

## Resume Training from Checkpoint

To resume training from a saved checkpoint, you need to load the checkpoint and restore the models state, optimizer state, and other parameters.

```
python
# Load the checkpoint
checkpoint = torch.load(checkpoint_epoch_5.pth)

# Instantiate the model and optimizer
model = ...
optimizer = optim.SGD(model.parameters(), lr=0.001)

# Restore the model and optimizer states
model.load_state_dict(checkpoint[state_dict])
optimizer.load_state_dict(checkpoint[optimizer])
start_epoch = checkpoint[epoch]
best_loss = checkpoint[loss]

# Resume training loop
num_epochs = 10
for epoch in range(start_epoch, num_epochs):
    # Training code here
    # ...

    # Save model checkpoint (optional)
    checkpoint = {
        epoch: epoch + 1,
        state_dict: model.state_dict(),
        optimizer: optimizer.state_dict(),
        loss: loss,
        # Other training parameters
    }
    torch.save(checkpoint, f'checkpoint_epoch_{epoch + 1}.pth')
```

By following these steps, you can save model checkpoints during training and resume training using the saved checkpoints. This approach allows you to keep track of the models progress, maintain training continuity, and fine-tune the model as needed.

# Customizing Loss Functions and Metrics

Customizing loss functions and metrics is crucial when you're working on specific tasks that require tailored optimization objectives and evaluation criteria. PyTorch allows you to define your own loss functions and compute custom metrics based on your problem domain. Here's how:

## 1. Custom Loss Function:

Define a custom loss function by subclassing `torch.nn.Module` and implementing the `forward()` method.

```
python
import torch
import torch.nn as nn

class CustomLoss(nn.Module):
    def __init__(self, weight=None):
        super(CustomLoss, self).__init__()
        self.weight = weight

    def forward(self, predicted, target):
        # Implement your custom loss computation here
        # Calculate loss based on predicted and target values
        loss = torch.mean((predicted - target) ** 2)
        return loss
```

## 2. Custom Metric:

You can define custom evaluation metrics using standard PyTorch tensor operations.

```
python
def custom_accuracy(predicted, target):
    correct = (predicted == target).sum()
    total = target.size(0)
    accuracy = correct.float() / total
    return accuracy.item()
```

## Using TensorBoard for Visualization

TensorBoard is a powerful visualization tool provided by TensorFlow that can be used with PyTorch as well. It allows you to track and visualize various aspects of your models training and performance over time.

### 1. Install TensorBoard:

Install the `tensorboard` library using pip:

```
bash
pip install tensorboard
```

### 2. Import Necessary Libraries:

```
python
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.tensorboard import SummaryWriter
```

### 3. Initialize TensorBoard Writer:

```
python
writer = SummaryWriter(logs) # Provide a directory path for log storage
```

### 4. Inside Training Loop:

Inside your training loop, add code to log loss, custom metrics, and other relevant information to TensorBoard.

```
python
for epoch in range(num_epochs):
    model.train()
    optimizer.zero_grad()
    outputs = model(x)
    loss = criterion(outputs, y)
    loss.backward()
    optimizer.step()

    # Log loss to TensorBoard
    writer.add_scalar(Loss/train, loss.item(), epoch)

    # Calculate and log custom metric
    accuracy = custom_accuracy(outputs.argmax(dim=1), y)
    writer.add_scalar(Accuracy/train, accuracy, epoch)
```

### 5. Run TensorBoard:

Open a terminal and navigate to the directory where you stored your logs. Run the following command to start TensorBoard:

```
bash  
tensorboard --logdir=logs
```

This will launch TensorBoard in your web browser, where you can visualize various metrics, loss curves, and other information over the course of training.

By customizing loss functions, metrics, and using tools like TensorBoard, you can tailor your deep learning process to your specific tasks and gain insights into your models performance in a more detailed and informative manner.